

SUSILKESSAV SESHADRI BHUVANESWARI

Boston, MA — +1 (617) 396-1430 — seshadribhuvanewa.s@northeastern.edu — [linkedin.com/in/susilkessav](https://www.linkedin.com/in/susilkessav) — github.com/Susilkessav

SKILLS

Languages : Python, SQL, TypeScript.
Data Engineering : PySpark, Airflow, ETL/ELT, dbt, Delta Lake, Great Expectations.
AI / LLM : LangChain, RAG, Ollama, ChromaDB, OpenAI API, LoRA, PEFT, Langfuse.
Machine Learning : TensorFlow, scikit-learn, ONNX, Computer Vision, Anomaly Detection.
Cloud / MLOps : AWS, Azure Synapse, Databricks, GCP, Docker, GitHub Actions, CI/CD.
Databases : PostgreSQL, MySQL, MongoDB, SAP HANA.

WORK EXPERIENCE

Data Engineer Intern **Jul 2024 – Sep 2025**
Getinge Wayne, United States

- Designed 15+ **ELT** ingestion pipelines on Azure Synapse and Databricks, pulling **50M+** daily records at 99.9% SLA.
- Architected a medallion Bronze Silver Gold pipeline with **contracts** and schema validated Parquet, standardizing **100+** tables for analytics.
- Built **GitHub Actions** CI/CD with schema validation gates and automated rollback across **15+** pipelines, enabling zero downtime deployments.
- Optimized 15+ **Airflow** DAGs with task parallelism and Dockerized isolation, cutting nightly batch runtime **28%** across the ingestion fleet.
- Deployed **Great Expectations** schema-drift detection across **100+** tables, cutting data triage time 20% by catching contract breaks early.
- Reduced service failures **40%** across 12 production workflows using idempotent retries, contract checks, and automated alerting thresholds.

Software Developer Intern **Aug 2022 – Feb 2023**
CISAI Amritapuri, India

- Built an automated **Python** collection service streaming **10,000+** posts daily into an Airflow ELT layer feeding NLP pipelines.
- Modeled **anomaly detection** over high-volume streaming logs, improving early outlier identification **25%** and cutting false-positive alert noise.
- Fine tuned LLM embedding models using **LoRA** and PEFT on domain datasets, boosting NLP accuracy **10%** with efficient training.
- Containerized inference microservices with **Docker** and deployed to AWS, cutting deployment time **30%** by removing drift across releases.

AI Engineer Intern **Dec 2021 – Jun 2022**
EasyCrop Tech Hyderabad, India

- Shipped a model serving pipeline on **AWS SageMaker**, processing 1,000+ images daily at **sub-150ms** P95 latency via REST endpoints.
- Automated model releases via **S3** event triggers, Lambda functions, and SageMaker pipelines, saving **8+** hours per deployment cycle.
- Configured **GitHub Actions** pipelines with Docker and **pytest**, reaching 88% coverage across unit, integration, and end-to-end suites.
- Applied **ONNX** quantization and transfer learning, compressing model size **40%** while preserving 95% accuracy on the SageMaker endpoint.

PROJECTS

ValoCoach – Local-First Agentic AI Coaching CLI — [GitHub](#)

- Architected a local-first agentic AI CLI using **Ollama** and **RAG**, routing 9 intent types through controlled elicitation and guardrails.
- Grounded **LLM** responses on a **342-doc** ChromaDB base spanning 22 agents, with deterministic fact-checking and 3-question elicitation caps.
- Engineered a **10 analyzer** evaluation harness scoring tactical decisions across a **20 turn** REPL window for repeatable feedback.
- Persisted match history in **SQLite** with Alembic migrations and validated agent workflows with **pytest** and continuous integration runs.

Amazon Automotive MLOps – RAG Data Platform — [GitHub](#)

- Engineered an end to end **LangChain** RAG system combining Airflow ingestion, **ChromaDB** retrieval, SentenceTransformer embeddings.
- Instrumented **Langfuse** observability across **100%** of retrieval calls, tracking latency, retrieval drift, and LLM output quality metrics.
- Designed a Dockerized **GCP** deployment blueprint with monitoring hooks and **OpenAI** moderation safeguards for the serving layer.

AI Sports Commentary Engine — [GitHub](#)

- Created a real time LLM platform with a **4-service** microarchitecture spanning ingestion, streaming enrichment, generation, and serving.
- Processed live play events through **Kafka** and Redpanda topics with **WebSockets**, delivering low-latency updates to the unified dashboard.
- Modeled live win-probability and hype-scoring logic conditioning **LLM** commentary tone on game state across high-impact and routine plays.

PUBLICATIONS

Peer-to-Peer Energy Trading Using Blockchain in Microgrid — IEEE (2023)

EDUCATION

Northeastern University **Sep 2023 – Dec 2025**
Master of Science in Data Analytics Engineering Boston, Massachusetts

Coursework: Algorithms, Distributed Systems, Database Management, Operating Systems, Machine Learning

Amrita Vishwa Vidyapeetham **Aug 2019 – May 2023**
B.Tech in Electrical and Computer Engineering Kerala, India

Coursework: Data Structures, Computer Architecture, Operating Systems, Object-Oriented Programming, Linear Algebra